


OS Support for Building Distributed Applications: Multithreaded Programming using Java Threads



Dr. Rajkumar Buyya

Cloud Computing and **D**istributed **S**ystems (CLOUDS) Laboratory
School of Computing and Information Systems
The University of Melbourne, Australia


<http://www.buyya.com>

Outline

- Introduction to Middleware
- Thread Applications
- Defining Threads
- Java Threads and States
- Architecture of Multithreaded servers
- Threads Synchronization
- Summary

Introduction

- Middleware is a layer of software (system) between Applications and Operating System (OS) powering the nodes of a distributed system.
- The OS facilitates:
 - Encapsulation and protection of resources inside servers;
 - Invocation of mechanisms required to access those resources including concurrent access/processing.




Middleware and Network Operating System (NOS)

- Many DOS (Distributed OS) have been investigated, but there are none in general/wide use. But NOS are in wide use for various reasons both technical and non-technical.
 - Users have much invested in their application software; they will not adopt a new OS that will not run their applications.
 - Users tend to prefer to have a degree of autonomy of their machines, even in a closely knit organisation.
- A combination of middleware and NOSs provides an acceptable balance between the requirement of autonomy and network transparency.
 - NOS allows users to run their favorite word processor.
 - Middleware enables users to take advantage of services that become available in their distributed systems.

Introducing a middleware

■ Building Distributed Systems

- DOS or NOS are not enough to build a DS!
- NOS are a good starting point but
- ... we need an additional layer “gluing” all together



Building Distributed Systems

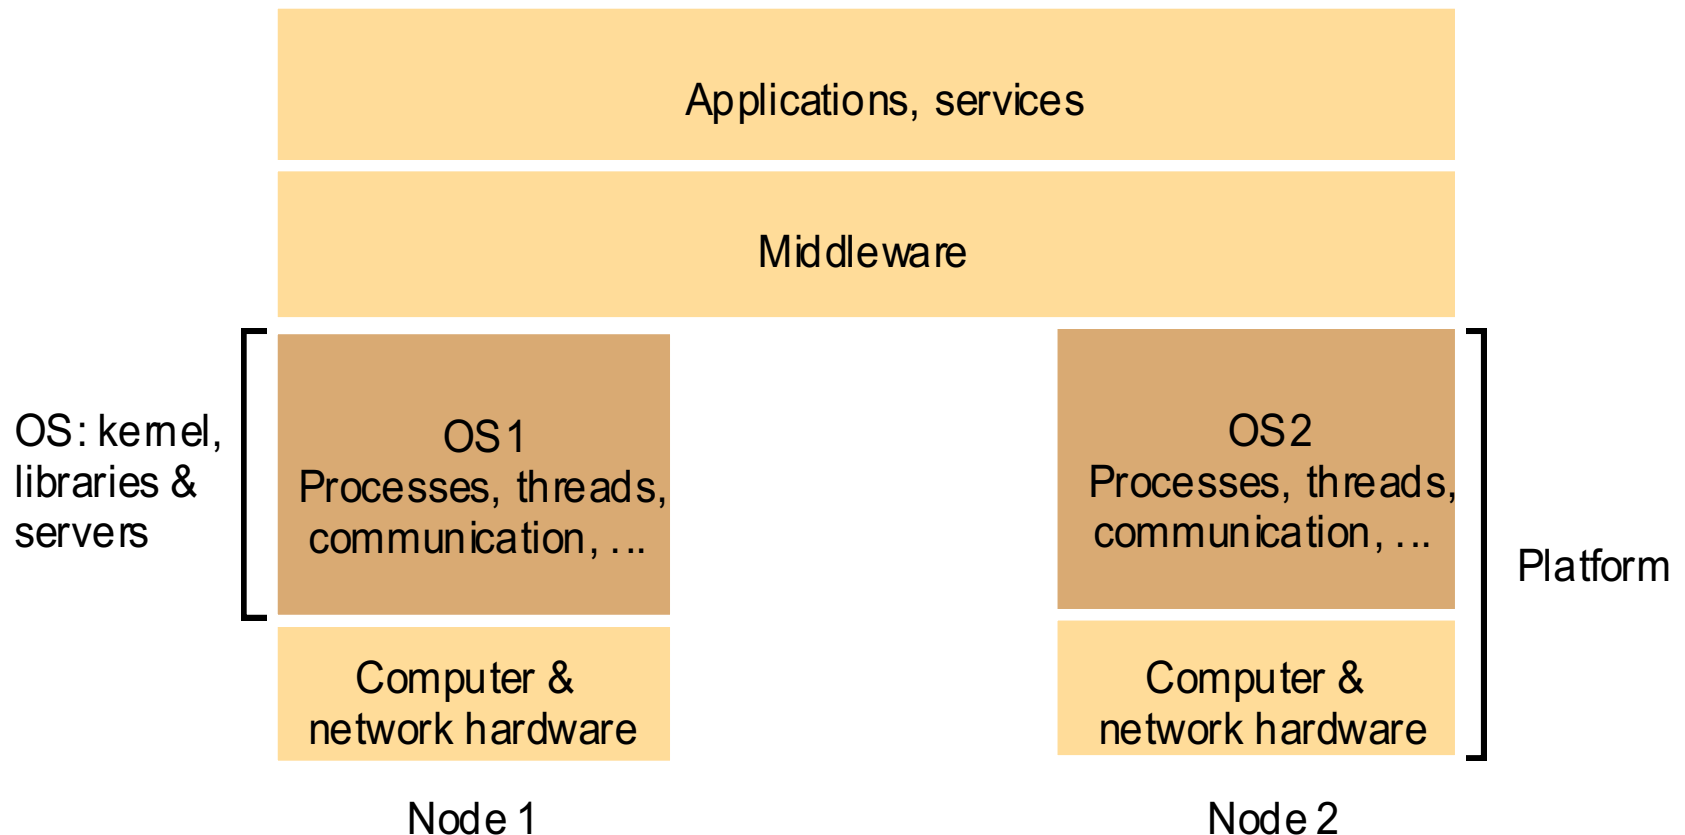
■ Middleware

- High-level features for DS
 - Communication
 - Management
 - Application specific
- Uniform layer where to build DS services
- Runtime environment of applications

■ Operating System

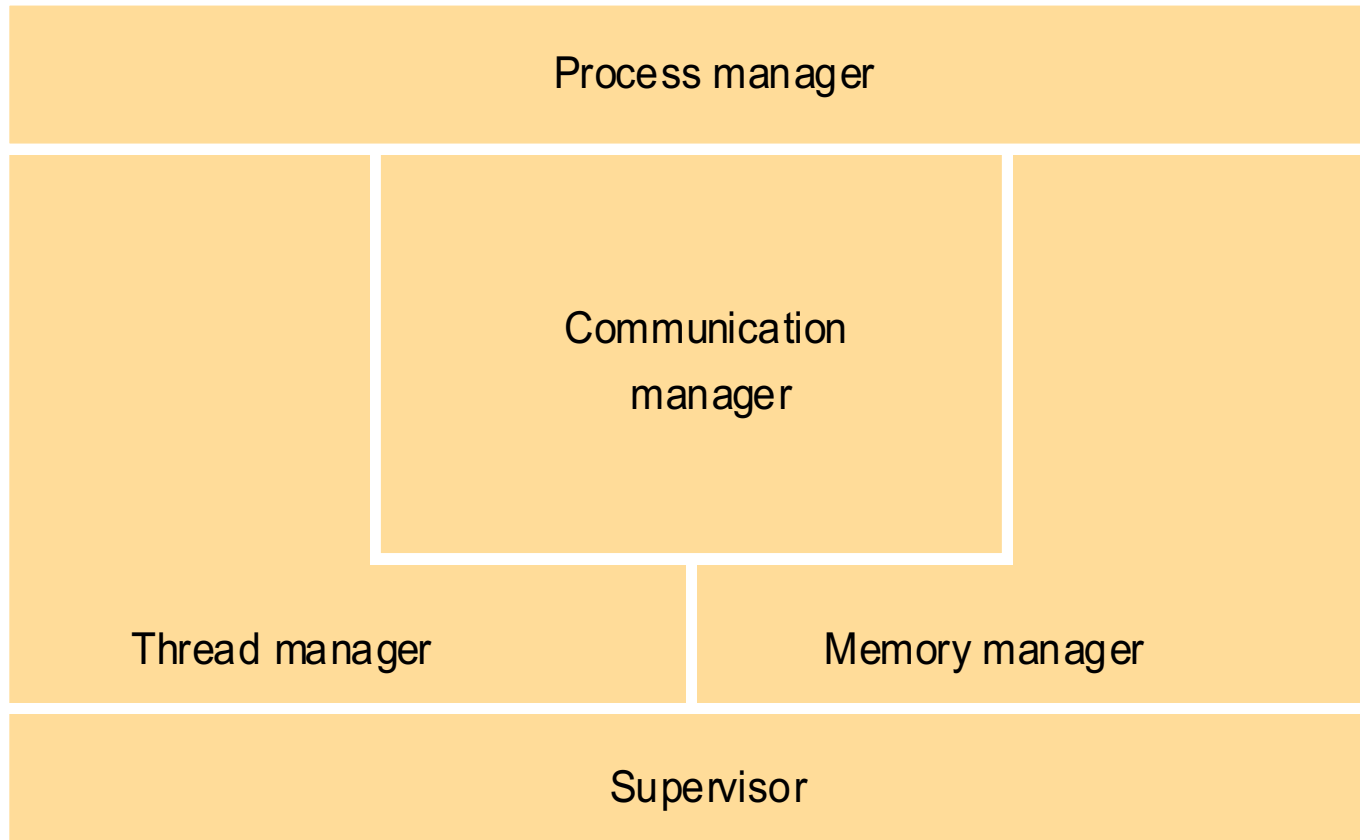
- Low / medium level (core) features
 - Process / threads management
 - Local hardware (CPU, disk, memory)
 - Security (users, groups, domain, ACLs)
 - Basic networking

Operating system layers and Middleware



- Unix and Windows are two examples of Network Operating Systems – have a networking capability built into them and so can be used to access remote resources using basic services such as rlogin, telnet.

Core OS components and functionality



Threaded Applications

■ Modern Applications and Systems

■ Operating System Level

- Multitasking: multiple applications running at once

■ Application Level

- Multithreading: multiple operations performed at the same time within an application.

■ Bottom Line:

- Illusion of concurrency

Threaded Applications

■ Modern Systems

- Multiple applications run concurrently!
- This means that... there are multiple processes on your computer



A single threaded program

```
class ABC
```

```
{
```

```
....
```

```
    public void main(..)
```

```
    {
```

```
    ...
```

```
    ..
```

```
    }
```

```
}
```

begin

body

end




Threaded Applications

■ Modern Systems

- Applications perform many tasks at once!
- This means that... there are multiple threads within a single process.




A Multithreaded Program



Threads may switch or exchange data/results

Single and Multithreaded Processes


threads are light-weight processes within a process



Multithreaded Server: For Serving Multiple Clients Concurrently

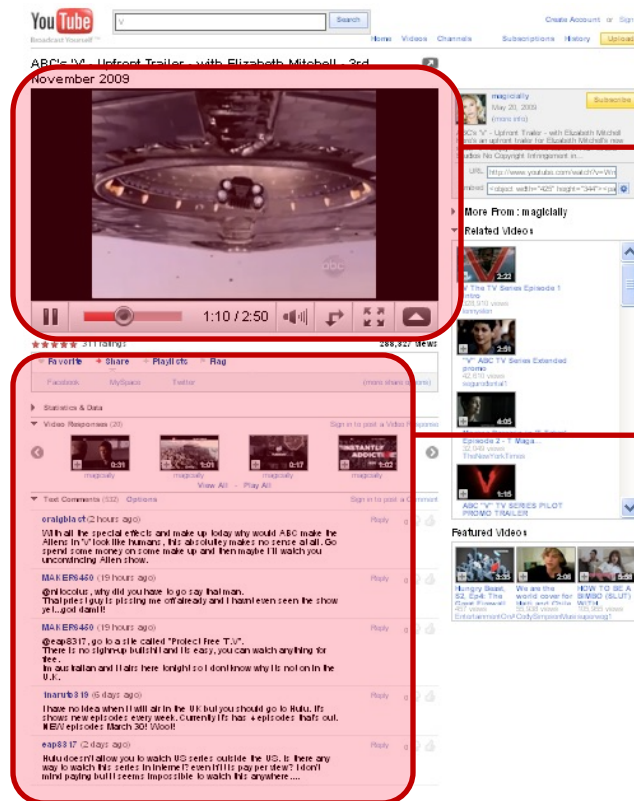
- Modern Applications

- Example: Multithreaded Web Server



Threaded Applications


- Modern Applications
 - Example: Internet Browser + Youtube



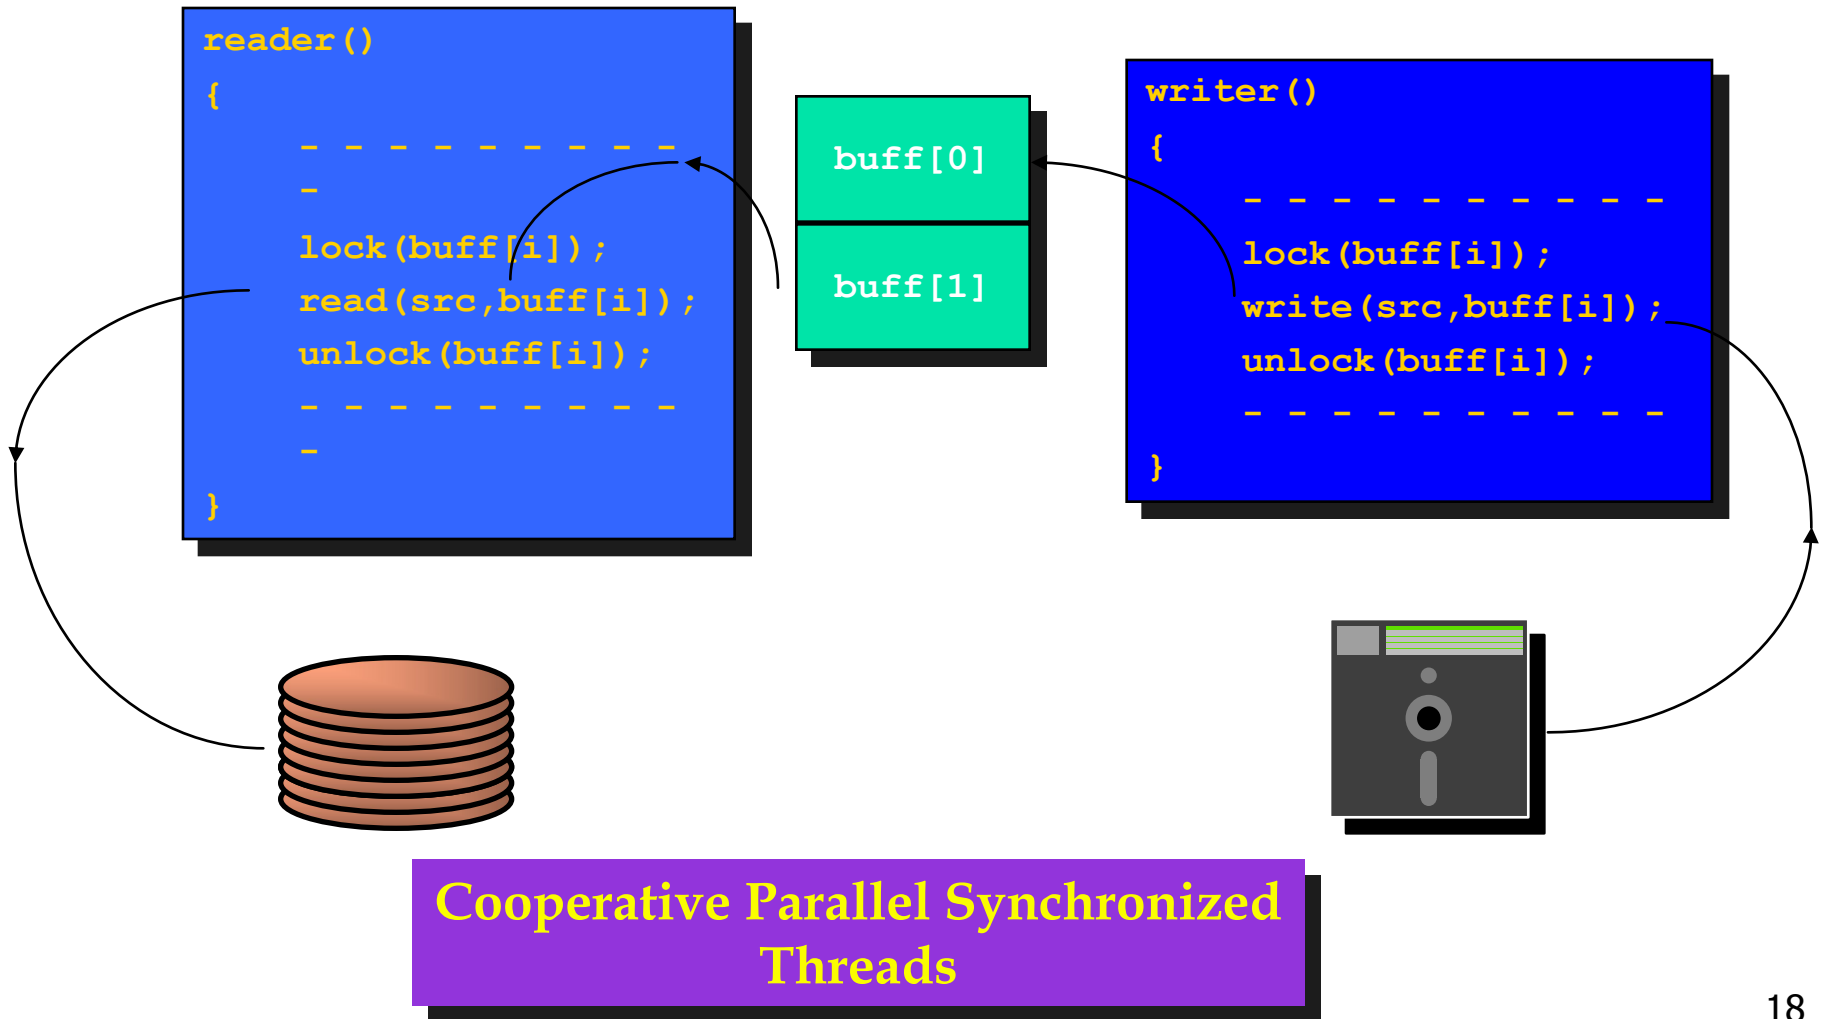
Video Streaming

Favorites, Share, Comments Posting

Modern Applications need Threads (ex1): Editing and Printing documents in background.



Multithreaded/Parallel File Copy

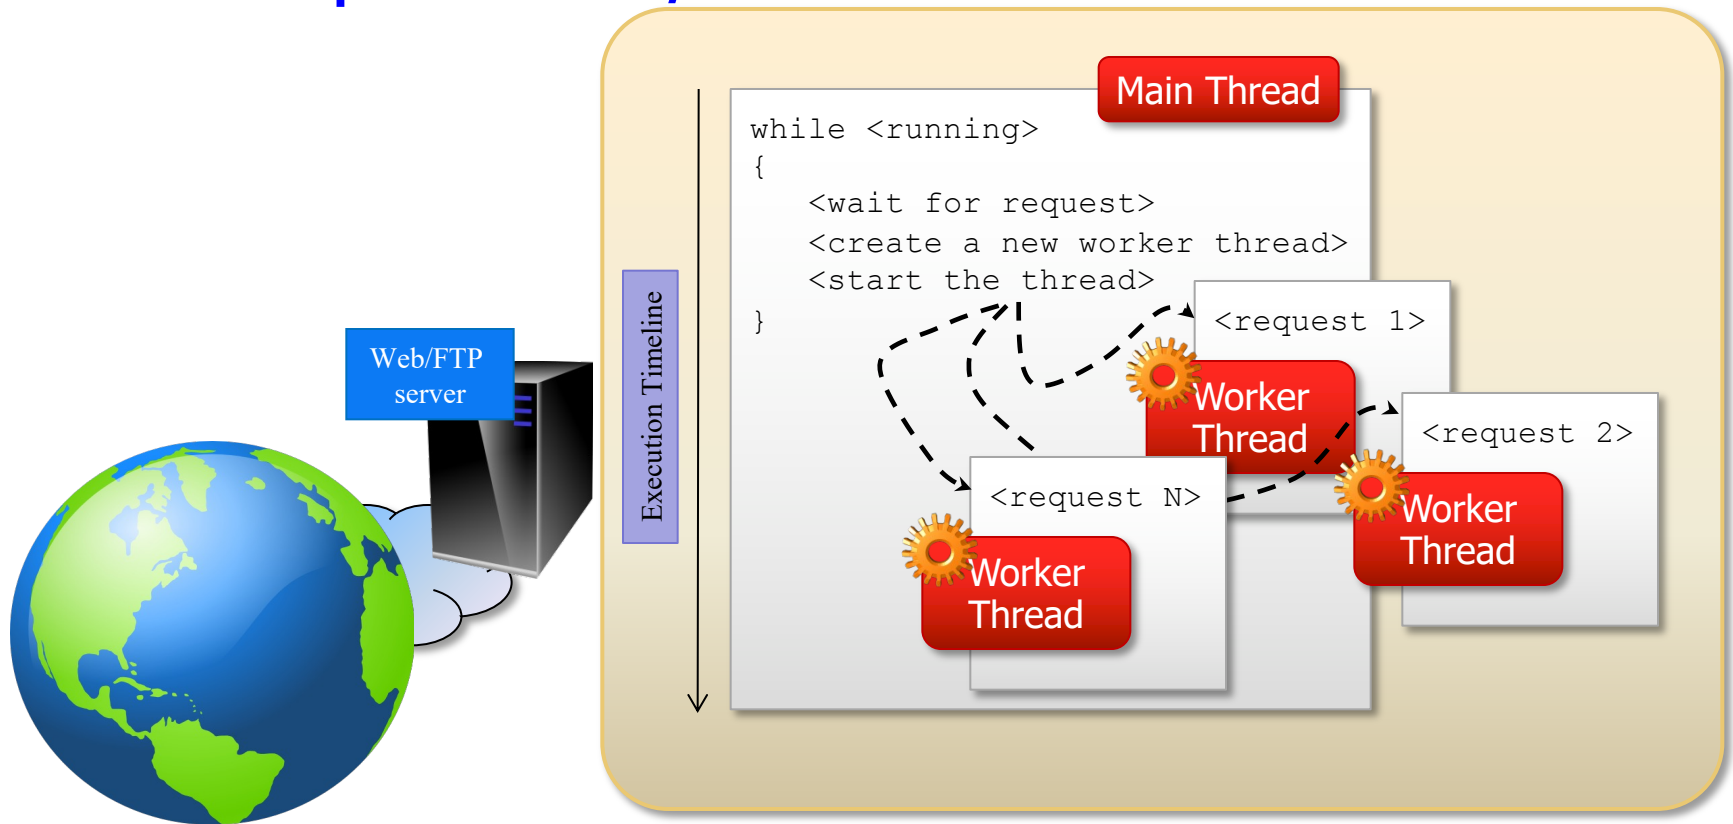


Defining Threads

- Applications – Threads are used to perform:
 - Parallelism and concurrent execution of independent tasks / operations.
 - Implementation of reactive user interfaces.
 - Non blocking I/O operations.
 - Asynchronous behavior.
 - Timer and alarms implementation.

Defining Threads

- Example: Web/FTP Server



Defining Threads


- A Thread is a piece of code that runs in **concurrent** with other threads.
- Each thread is a statically ordered sequence of instructions.
- Threads are used to **express concurrency** on both single and multiprocessors machines.
- Programming a task having multiple threads of control is called: Multithreading or Multithreaded Programming.

Java Threads


- Java has built in support for Multithreading
- Synchronization
- Thread Scheduling
- Inter-Thread Communication:
 - | | | |
|---------------|-------|-------------|
| currentThread | start | setPriority |
| yield | run | getPriority |
| sleep | stop | suspend |
| resume | | |
- Java Garbage Collector is a low-priority thread.

Threading Mechanisms...

- Create a class that extends the Thread class
- Create a class that implements the Runnable interface



[a]



[b]

1st method: Extending Thread class

- Create a class by extending Thread class and override run() method:

```
class MyThread extends Thread
{
    public void run()
    {
        // thread body of execution
    }
}
```

- Create a thread:

```
MyThread thr1 = new MyThread();
```

- Start Execution of threads:

```
thr1.start();
```

- Create and Execute:

```
new MyThread().start();
```


An example

```
class MyThread extends Thread {  
    public void run() {  
        System.out.println(" this thread is running ... ");  
    }  
}
```

```
class ThreadEx1 {  
    public static void main(String [] args ) {  
        MyThread t = new MyThread();  
        t.start();  
    }  
}
```

2nd method: Threads by implementing Runnable interface

- Create a class that implements the interface Runnable and override run() method:

```
class MyThread implements Runnable
{
    .....
    public void run()
    {
        // thread body of execution
    }
}
```

- Creating Object:

```
MyThread myObject = new MyThread();
```

- Creating Thread Object:

```
Thread thr1 = new Thread( myObject );
```

- Start Execution:

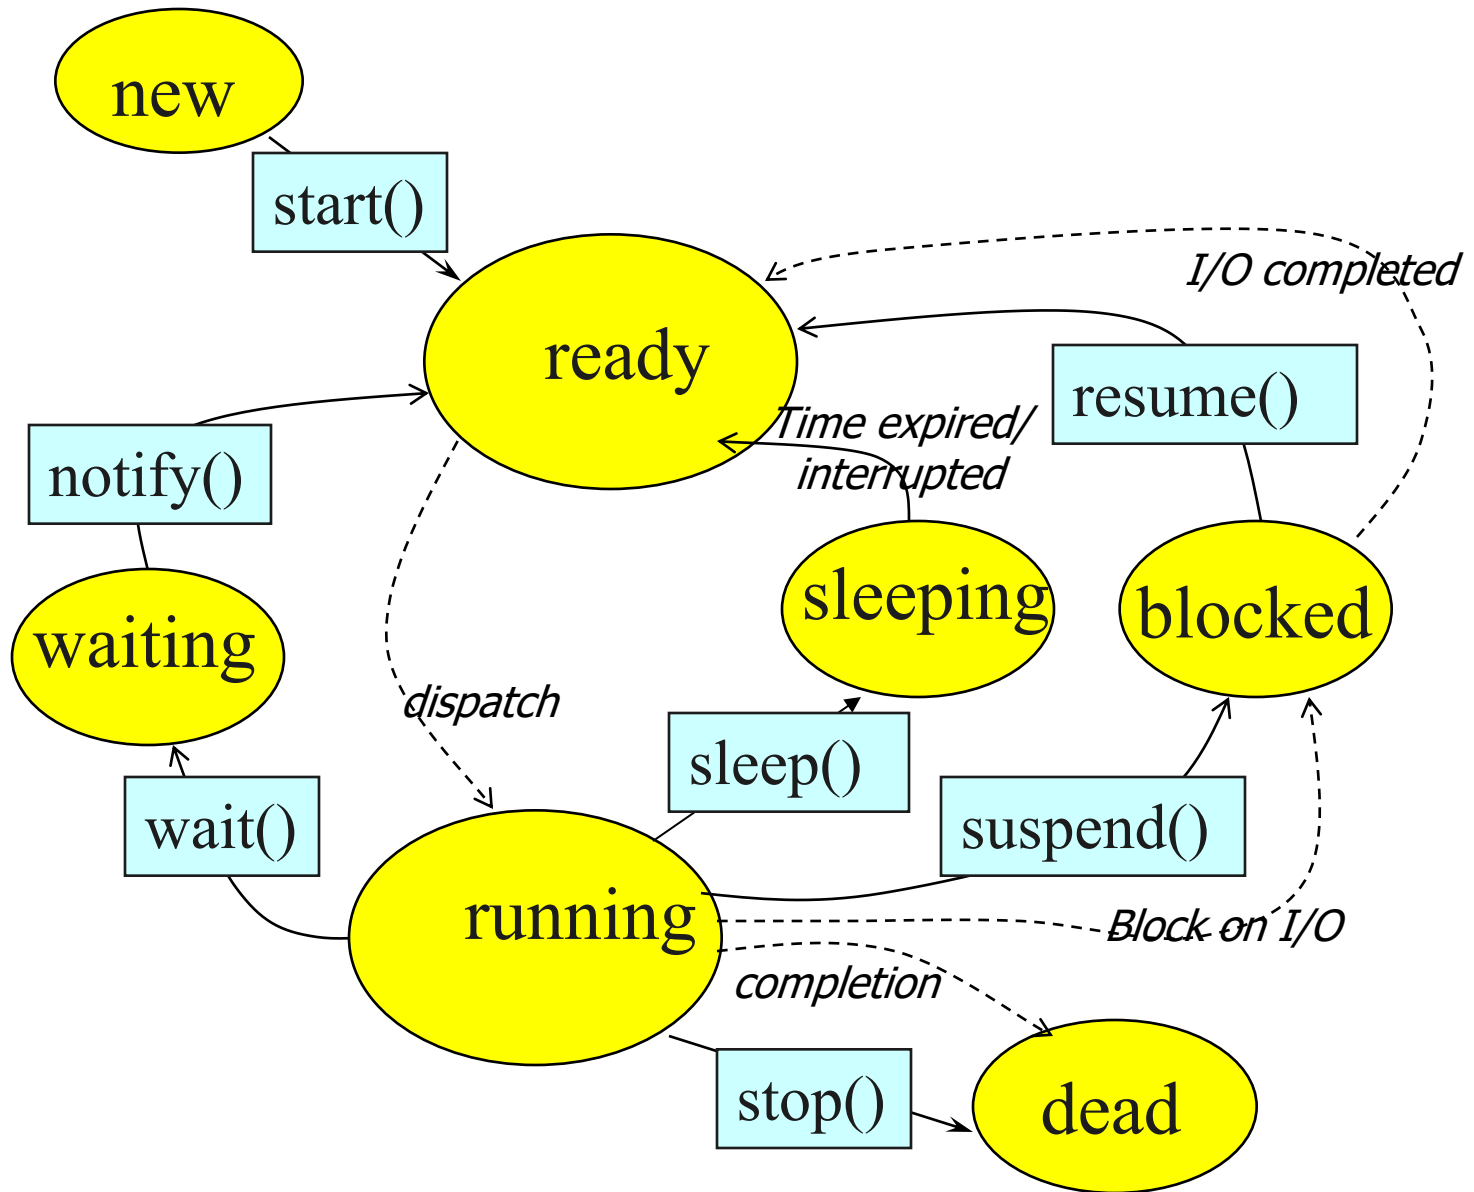
```
thr1.start();
```

An example

```
class MyThread implements Runnable {  
    public void run() {  
        System.out.println(" this thread is running ... ");  
    }  
}
```

```
class ThreadEx2 {  
    public static void main(String [] args ) {  
        Thread t = new Thread(new MyThread());  
        t.start();  
    }  
}
```

Life Cycle of Thread



A Program with Three Java Threads

- Write a program that creates 3 threads

Three threads example

```
■ class A extends Thread
■ {
■     public void run()
■     {
■         for(int i=1;i<=5;i++)
■         {
■             System.out.println("\t From ThreadA: i= "+i);
■         }
■         System.out.println("Exit from A");
■     }
■ }

■ class B extends Thread
■ {
■     public void run()
■     {
■         for(int j=1;j<=5;j++)
■         {
■             System.out.println("\t From ThreadB: j= "+j);
■         }
■         System.out.println("Exit from B");
■     }
■ }
```

Three threads example

```
■ class C extends Thread
■ {
■     public void run()
■     {
■         for(int k=1;k<=5;k++)
■         {
■             System.out.println("\t From ThreadC: k= "+k);
■         }
■
■         System.out.println("Exit from C");
■     }
■ }

■ class ThreadTest
■ {
■     public static void main(String args[])
■     {
■         new A().start();
■         new B().start();
■         new C().start();
■     }
■ }
```

Run 1

- [raj@mundroo] threads [1:76] java ThreadTest

From ThreadA: i= 1

From ThreadA: i= 2

From ThreadA: i= 3

From ThreadA: i= 4

From ThreadA: i= 5

Exit from A

From ThreadC: k= 1

From ThreadC: k= 2

From ThreadC: k= 3

From ThreadC: k= 4

From ThreadC: k= 5

Exit from C

From ThreadB: j= 1

From ThreadB: j= 2

From ThreadB: j= 3

From ThreadB: j= 4

From ThreadB: j= 5

Exit from B

Run 2

- [raj@mundroo] threads [1:77] java ThreadTest

From ThreadA: i= 1

From ThreadA: i= 2

From ThreadA: i= 3

From ThreadA: i= 4

From ThreadA: i= 5

From ThreadC: k= 1

From ThreadC: k= 2

From ThreadC: k= 3

From ThreadC: k= 4

From ThreadC: k= 5

Exit from C

From ThreadB: j= 1

From ThreadB: j= 2

From ThreadB: j= 3

From ThreadB: j= 4

From ThreadB: j= 5

Exit from B

Exit from A

Thread Priority

- In Java, each thread is assigned priority, which affects the order in which it is scheduled for running. The threads so far had same default priority (NORM_PRIORITY) and they are served using FCFS policy.
 - Java allows users to change priority:
 - ThreadName.setPriority(intNumber)
 - MIN_PRIORITY = 1
 - NORM_PRIORITY=5
 - MAX_PRIORITY=10

Thread Priority Example

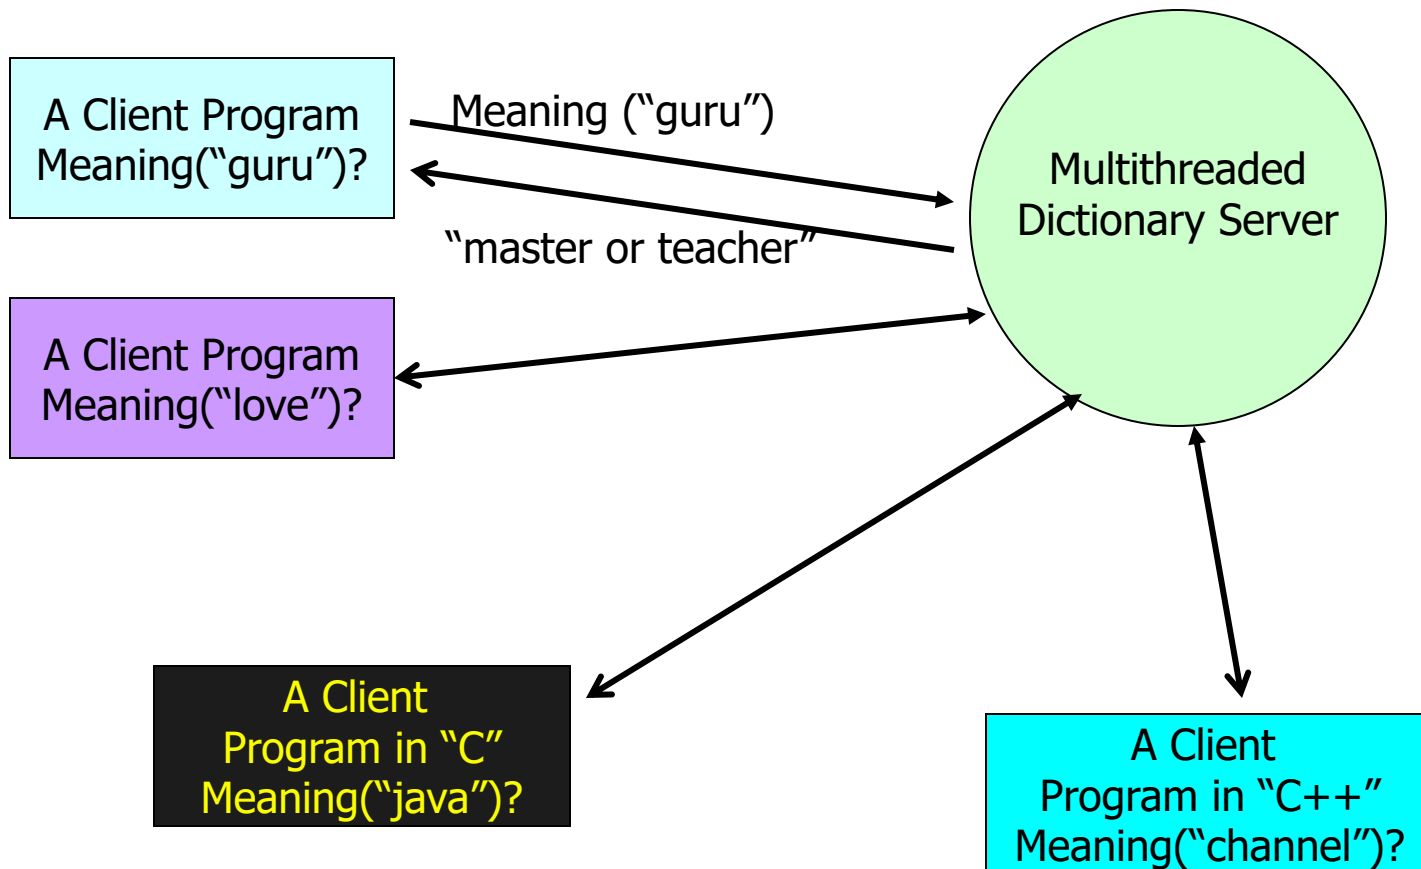
```
class A extends Thread
{
    public void run()
    {
        System.out.println("Thread A started");
        for(int i=1;i<=4;i++)
        {
            System.out.println("\t From ThreadA: i= "+i);
        }
        System.out.println("Exit from A");
    }
}

class B extends Thread
{
    public void run()
    {
        System.out.println("Thread B started");
        for(int j=1;j<=4;j++)
        {
            System.out.println("\t From ThreadB: j= "+j);
        }
        System.out.println("Exit from B");
    }
}
```

Thread Priority Example

```
class C extends Thread
{
    public void run()
    {
        System.out.println("Thread C started");
        for(int k=1;k<=4;k++)
        {
            System.out.println("\t From ThreadC: k= "+k);
        }
        System.out.println("Exit from C");
    }
}
class ThreadPriority
{
    public static void main(String args[])
    {
        A threadA=new A();
        B threadB=new B();
        C threadC=new C();
        threadC.setPriority(Thread.MAX_PRIORITY);
        threadB.setPriority(threadA.getPriority()+1);
        threadA.setPriority(Thread.MIN_PRIORITY);
        System.out.println("Started Thread A");
        threadA.start();
        System.out.println("Started Thread B");
        threadB.start();
        System.out.println("Started Thread C");
        threadC.start();
        System.out.println("End of main thread");
    }
}
```


Assignment 1 at a Glance: Multithreaded Dictionary Server – Using Sockets and Threads



Accessing Shared Resources

- Applications access to shared resources need to be coordinated.
 - Printer (two person jobs cannot be printed at the same time)
 - Simultaneous operations on your bank account.
 - Can the following operations be done at the same time on the same account?
 - Deposit()
 - Withdraw()
 - Enquire()

Online Bank: Serving Many Customers and Operations



Shared Resources



- If one thread tries to read the data and other thread tries to update the same data, it leads to inconsistent state.
- This can be prevented by synchronising access to the data.
- Use “synchronized” method:
 - `public synchronized void update()`
 - `{`
 - `...`
 - `}`

the driver: 3 Threads sharing the same object


```
class InternetBankingSystem {  
    public static void main(String [] args ) {  
        Account accountObject = new Account ();  
        Thread t1 = new Thread(new MyThread(accountObject));  
        Thread t2 = new Thread(new YourThread(accountObject));  
        Thread t3 = new Thread(new HerThread(accountObject));  
        t1.start();  
        t2.start();  
        t3.start();  
        // DO some other operation  
    } // end main()  
}
```

Shared account object between 3 threads

```
class MyThread implements Runnable {  
    Account account;  
    public MyThread (Account s) { account = s;}  
    public void run() { account.deposit(); }  
} // end class MyThread
```

```
class YourThread implements Runnable {  
    Account account;  
    public YourThread (Account s) { account = s;}  
    public void run() { account.withdraw(); }  
} // end class YourThread
```

```
class HerThread implements Runnable {  
    Account account;  
    public HerThread (Account s) { account = s; }  
    public void run() { account.enquire(); }  
} // end class HerThread
```



Monitor (shared object access): serializes operation on shared objects

```
class Account { // the 'monitor'
    int balance;


    // if 'synchronized' is removed, the outcome is unpredictable
    public synchronized void deposit( ) {
        // METHOD BODY : balance += deposit_amount;
    }

    public synchronized void withdraw( ) {
        // METHOD BODY: balance -= deposit_amount;
    }
    public synchronized void enquire( ) {
        // METHOD BODY: display balance.
    }
}
```

Architecture for Multithread Servers


- Multithreading enables servers to maximize their throughput, measured as the number of requests processed per second.
- Threads may need to treat requests with varying priorities:
 - A corporate server could prioritize request processing according to class of customers.
- Architectures:
 - Worker pool
 - Thread-per-request
 - Thread-per-connection
 - Thread-per-object

Client and server with threads (worker-pool architecture)




- In worker-pool architectures, the server creates a **fixed pool of worker threads** to process requests.
- The module "receipt and queuing" receives requests from sockets/ports and places them on a shared request queue for retrieval by the workers.

Alternative server threading architectures




a. Thread-per-request

IO Thread creates a new worker thread for each request and worker thread destroys itself after serving the request.



b. Thread-per-connection

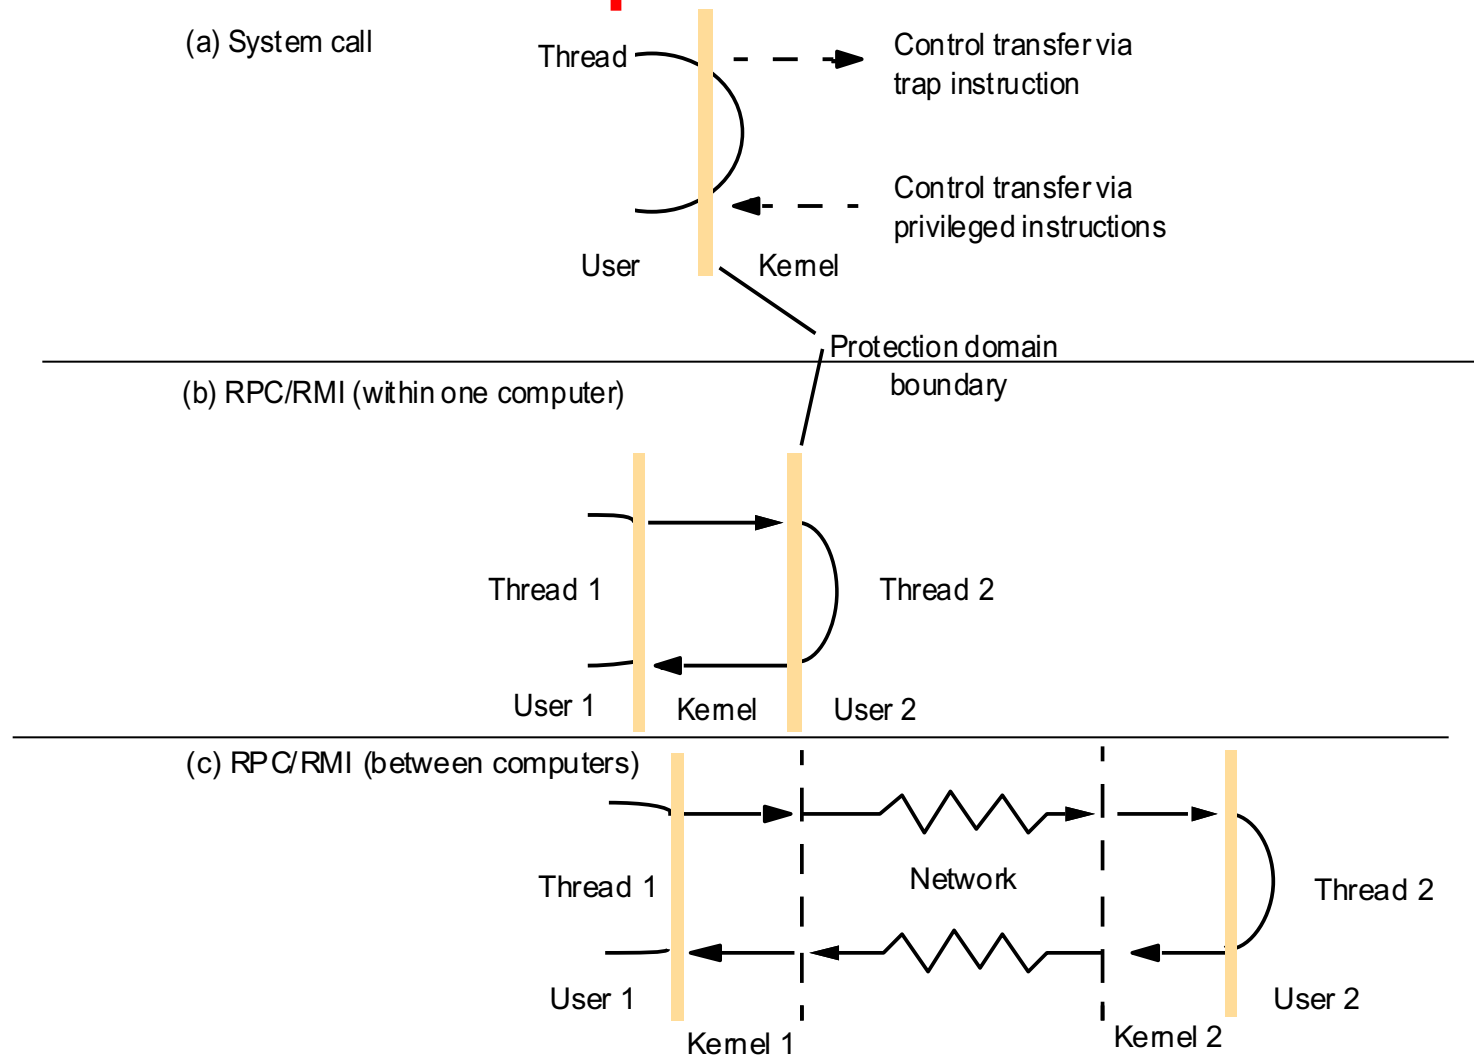
Server associates a Thread with each connection and destroys when client closes the connection. Client may make many requests over the connection.



c. Thread-per-object

Associates Thread with each object. An IO thread receives request and queues them for workers, but this time there is a **per-object queue**.

Invocations between address spaces



Summary

- Operating system provides various types of facilities to support middleware for distributed system:
 - encapsulation, protection, and concurrent access and management of node resources.
- Multithreading enables servers to maximize their throughput, measured as the number of requests processed per second.
- Threads support treating of requests with varying priorities.
- Various types of architectures can be used in concurrent processing:
 - Worker pool
 - Thread-per-request
 - Thread-per-connection
 - Thread-per-object
- Threads need to be synchronized when accessing and manipulating shared resources.
- New OS designs provide flexibility in terms of separating mechanisms from policies.

References

- CDK Book (Text Book)
 - Chapter 7 – “Operating System Support”
- Chapter 14: Multithread Programming
 - R. Buyya, S. Selvi, X. Chu, **“Object Oriented Programming with Java: Essentials and Applications”**, McGraw Hill, New Delhi, India, 2009.